

CIS3110 Lecture 9 - Summary Notes: Disk Structure and File Systems

1. Disk Structure and Addressing (Chapter 11.1, 11.2)

Three-Dimensional Disk Addressing

- Disks use a three-dimensional addressing scheme:
 - **Theta**: Angular position (sectors around the disk)
 - **R**: Radius (tracks from inner to outer edge)
 - **Z**: Head position (which platter surface)
- This forms a "cubical" addressing space despite the circular physical shape
- Standard sector size: 512 bytes (hardware-defined unit)

Physical Disk Components

- Hard drives contain multiple platters with read/write heads
- Heads are positioned by movable arms over the spinning platters
- Modern drives: multiple platters in sealed cases
- Historical perspective: early drives used removable "cartridges" with platters

2. Disk Scheduling Algorithms (Chapter 12.4)

Problem: Multiple Disk Requests

- Moving the disk arm is costly in terms of time
- Goal: Minimize head movement while serving all requests

Algorithms

1. FIFO (First-In, First-Out)

- Serve requests in order of arrival
- Problem: Inefficient movement, excessive head travel

2. Shortest Job First (Shortest Seek Time First)

- Always go to the closest request
- Problem: Starvation for distant requests

3. Elevator/SCAN Algorithm

- Similar to modern elevator operation

- Move arm in one direction until no more requests in that direction
- Then reverse direction and repeat
- Picks up any new requests that are "ahead" of current position
- Guarantees service to all requests eventually
- Better utilization of head movement

3. Block Size Considerations (Chapter 11.2, 11.3)

Block vs. Sector

- **Sector:** Hardware-defined unit (typically 512 bytes)
- **Block:** File system unit (comprised of multiple sectors)

Block Size Trade-offs

- **Larger blocks:**
 - More efficient transfers (more data per seek)
 - Worse internal fragmentation (wasted space for small files)
- **Smaller blocks:**
 - Less wasted space
 - More seeking and indexing overhead

4. FAT File System (Chapter 14.3, 14.4)

Structure

1. Master Boot Record (MBR)

- First block on disk
- Contains:
 - Jump location for boot code
 - Sector size information
 - Number of FAT copies
 - Disk geometry (sectors per track, heads, etc.)
 - Root directory entries count
 - Media type code
 - OEM name

2. File Allocation Table (FAT)

- Two identical copies for redundancy (protection against corruption)
- Acts as a linked list of block pointers
- Each entry corresponds to one logical block on disk
- Special values for "end of file" and "free block"

3. Root Directory

- Fixed-size array of directory entries (limited to 1024 entries)
- Each entry contains:
 - 8.3 format filename (8 chars + 3 char extension)
 - File attributes (read-only, hidden, etc.)
 - Modification time
 - File size in bytes
 - First block number

4. Data Blocks

- Rest of the disk stores actual file data
- Subdirectories are just special files with directory entries

FAT as a File Structure

- Files are represented as linked lists of blocks
- FAT table entries serve as "next pointers"
- Parallel structure: data blocks contain content, FAT contains linking information
- Disadvantages:
 - Poor random access (must traverse list from beginning)
 - Disk fragmentation issues

Historical Evolution of Block Sizes

- As disk sizes increased, Microsoft increased cluster (block) sizes
- For 16-bit FAT:
 - 360K disk: 1-2 sectors per cluster
 - 4GB disk: 64KB clusters
 - 16GB disk: 256KB clusters
- Larger cluster sizes lead to significant space wastage for small files
- NTFS follows similar pattern but supports much larger disks (up to 256TB)

5. Unix File System Overview (Chapter 15.5, 15.6)

Inode-Based Structure

- Uses a tree structure instead of linked list
- Root of file structure is the **inode** (index node)

Components

1. Super Block

- Contains file system parameters
- Multiple redundant copies throughout disk
- Includes:
 - First data block offset
 - File system size
 - Root inode location
 - Free space information
 - Device geometry
 - Dirty flag for caching

2. Inodes

- Complete descriptor for each file
- Contains:
 - File mode (permissions: read, write, execute)
 - Link count (number of directory entries pointing to this inode)
 - Owner/Group IDs
 - File size
 - Access/modification times
 - Direct block pointers (first 6 in example)
 - Indirect block pointers (for larger files)

3. File Structure (Tree-based)

- Small files: Direct block pointers in inode
- Medium files: First indirect pointer (addresses ~1024 more blocks)
- Large files: Second indirect pointer (addresses ~1024² more blocks)
- Very large files: Third indirect pointer (addresses ~1024³ more blocks)

Advantages over FAT

- Efficient random access (can directly access any block)
- Short, broad tree structure minimizes seek time
- Multiple hard links possible (same file appears in multiple directories)

6. SSD vs. Traditional Hard Drives

- SSDs don't have mechanical heads, so seek time isn't a factor
- Addressing structure remains similar, but physical layout differs
- Elevator/SCAN algorithms less relevant for SSDs

Next Lecture

- More details on Unix file system
- Comparison with NTFS
- Additional features of inodes (link count usage)